



MARMARA
UNIVERSITY

Embedded Digital Image Processing

EE4065

Homework 6

Mustafa Öztürk

150722016

Emirhan Şener

150721069

Faculty of Engineering
Electrical Electronic Engineering

Fall 2025

ACRONYMS

CNN Convolutional Neural Network

EE4065 Embedded Digital Image Processing

Hu Moments Hu Invariant Moments

MCU Microcontroller Unit

MNIST Modified National Institute of Standards and Technology

RAM Random Access Memory

UART Universal Asynchronous Receiver Transmitter

CONTENTS

Acronyms	1
List of Figures	3
1 Introduction	4
2 Problems	5
2.1 Q-1) Section 13.7 Application: Handwritten Digit Recognition from Digital Images.....	5
3 Conclusion	23
References	24

LIST OF FIGURES

2.1 ST Edge AI Developer Cloud analysis dashboard showing SqueezeNetV1.1 memory usage exceeds in flash size.	15
2.2 Successful generation of the STM32CubeIDE project for ResNetV1.....	16
2.3 Training accuracy and loss curves for all evaluated models.	20
2.4 Confusion matrices comparing the classification performance of the six architectures.	21
2.5 Real-time inference results from the STM32F446RE visualized via the Python client and showing computed probabilities of the label shown in STM32CubeIDE	22

1. INTRODUCTION

This report presents the implementation and results of the 6th homework for the Embedded Digital Image Processing (EE4065) course. Building upon previous assignments where we explored basic neural networks and feature extraction, this work focuses on deploying deep Convolutional Neural Network (CNN) architectures for image classification. The primary objective is to implement handwritten digit recognition on the STM32F446RE Microcontroller Unit (MCU), balancing model accuracy with the strict resource constraints of an embedded system.

The tasks in this assignment are based on Section 13.7 of the course textbook [1]. We utilized the Modified National Institute of Standards and Technology (MNIST) dataset to train and evaluate several standard architecture families, including EfficientNet, ResNet, MobileNet, and SqueezeNet. The goal was to compare these models not only on classification performance but also on their deployability regarding inference speed and memory footprint.

We used the **ST Edge AI Developer Cloud** [2] to optimize and convert our trained TensorFlow models into C code executable on the STM32. The workflow involved converting models from .h5 to .tflite format and applying INT8 quantization to minimize size. During this process, we observed that the optimized EfficientNet and SqueezeNet models exceeded the 512kB flash memory limit of the target MCU, rendering them unsuitable for this specific hardware despite sufficient Random Access Memory (RAM). Consequently, the on-board implementation phase focuses on ResNetV1, MobileNetV1, FD MobileNetV1, and ST FD MobileNetV1.

All models were initially trained and validated on a PC using Python and TensorFlow [3]. We analyzed the performance using accuracy metrics and confusion matrices. The compatible models were then generated as optimized projects via the cloud service and deployed to the MCU to demonstrate real-time inference capabilities.

2. PROBLEMS

2.1. Q-1) SECTION 13.7 APPLICATION: HANDWRITTEN DIGIT RECOGNITION FROM DIGITAL IMAGES

Implement the end-of-chapter application from the course textbook [1]. The goal is to perform Handwritten Digit Recognition from Digital Images using CNN models. You are required to use all the mentioned CNN models (SqueezeNet, EfficientNet, etc.) in the course and implement the solution on the target hardware.

Theory

CNNs and Quantization on Embedded Systems

Deploying CNNs on the STM32F446RE requires balancing model complexity with limited Flash (512kB) and RAM (128kB). A critical optimization is **Quantization**, which converts 32-bit floating-point weights to 8-bit integers. This reduces model size by 75% and accelerates inference on MCUs without floating-point units, utilizing integer arithmetic for matrix operations.

Efficient Architectures

We evaluated several architectures optimized for efficiency:

- **ResNetV1:** Utilizes residual skip connections to allow deeper networks while maintaining training stability.
- **MobileNetV1:** Designed for embedded vision, it uses *Depthwise Separable Convolutions* to drastically reduce parameters and computation compared to standard convolutions.

- **SqueezeNet & EfficientNet:** architectures that prioritize high accuracy-per-parameter. However, their complex graph structures often result in binary sizes that exceed the Flash memory of smaller MCUs like the F446RE.

ST Edge AI Developer Cloud

The **ST Edge AI Developer Cloud** serves as the deployment bridge. It optimizes the ‘.tflite’ models, validates memory constraints (Flash/RAM), and generates the optimized C library (X-CUBE-AI) required to run inference on the STM32 hardware.

Procedure

Step 1: Training the CNN Models

We trained a variety of CNN architectures to compare their performance and suitability for embedded deployment. The models include ResNetV1, SqueezeNetV1.1, EfficientNetV2-B0, MobileNetV1/V2, FDMobileNet, and ST-optimized variants.

Training Process: The models were trained using the MNIST dataset. Since MNIST images are 28x28 grayscale, we preprocessed them by resizing to 32x32 and converting them to 3-channel (RGB) format to match the input requirements of standard CNN backbones. Pixel values were normalized to the [0, 1] range.

Training was performed using the Adam optimizer with a learning rate of 0.001. We employed callbacks for model checkpointing (saving the best model based on validation accuracy), early stopping (patience of 5 epochs), and learning rate reduction on plateaus.

Listing 2.1: Training script for all CNN models (train_all_models.py)

```

1 import os
2 import sys
3 import numpy as np
4 import tensorflow as tf
5 from tensorflow import keras
6 import matplotlib
7 matplotlib.use('Agg')
8 import matplotlib.pyplot as plt
9 from sklearn.metrics import confusion_matrix, classification_report,
   accuracy_score, precision_score
10
11 sys.path.insert(0, os.path.join(os.path.dirname(__file__), 'supplementary'))
12
13 from resnetv1 import get_resnetv1
14 from squeezenetv11 import get_squeezenetv11

```

```

15 from efficientnetv2 import get_efficientnetv2
16 from mobilenetv1 import get_mobilenetv1
17 from mobilenetv2 import get_mobilenetv2
18 from fdmobilenet import get_fdmobilenet
19 from st_efficientnet_lc_v1 import get_st_efficientnet_lc_v1
20 from st_fdmobilenet_v1 import get_st_fdmobilenet_v1
21
22
23 # Configuration
24 NUM_CLASSES = 10
25 DATA_SHAPE = (32, 32, 3)
26 EPOCHS = 20
27 BATCH_SIZE = 64
28 PATIENCE = 5
29 MODELS_DIR = os.path.join(os.path.dirname(__file__), 'models')
30 RESULTS_DIR = os.path.join(os.path.dirname(__file__), 'results')
31
32
33 def prepare_tensor(images, out_shape):
34     images = tf.expand_dims(images, axis=-1)
35     images = tf.repeat(images, 3, axis=-1)
36     images = tf.image.resize(images, out_shape[:2])
37     images = images / 255.0
38     return images
39
40
41 def load_mnist_data():
42     print("Loading MNIST dataset...")
43     (train_images, train_labels), (val_images, val_labels) = tf.keras.
44         datasets.mnist.load_data()
45
46     print("Preprocessing data...")
47     train_images = prepare_tensor(train_images, DATA_SHAPE)
48     val_images = prepare_tensor(val_images, DATA_SHAPE)
49
50     train_labels = tf.keras.utils.to_categorical(train_labels,
51         NUM_CLASSES)
52     val_labels = tf.keras.utils.to_categorical(val_labels, NUM_CLASSES)
53
54     print(f"Training samples: {len(train_images)}")
55     print(f"Validation samples: {len(val_images)}")
56     print(f"Image shape: {train_images.shape[1:]}")
57
58     return (train_images, train_labels), (val_images, val_labels)
59
60
61 def create_callbacks(model_name):
62     model_path = os.path.join(MODELS_DIR, f'{model_name}.h5')
63
64     callbacks = [
65         keras.callbacks.ModelCheckpoint(
66             model_path,
67             save_best_only=True,
68             monitor='val_accuracy',
69             verbose=1
70         ),
71         keras.callbacks.EarlyStopping(
72             monitor='val_accuracy',
73             patience=PATIENCE,
74             verbose=1,
75             restore_best_weights=True
76         ),
77         keras.callbacks.ReduceLROnPlateau(
78             monitor='val_loss',
79         )
80     ]
81
82     return callbacks

```

```

77         factor=0.5,
78         patience=3,
79         verbose=1,
80         min_lr=1e-6
81     )
82 ]
83 return callbacks
84
85
86 def train_model(model, model_name, train_data, val_data):
87     print(f"\n{'='*60}")
88     print(f"Training {model_name}")
89     print(f"{'='*60}")
90
91     # Compile model
92     model.compile(
93         optimizer=keras.optimizers.Adam(learning_rate=1e-3),
94         loss='categorical_crossentropy',
95         metrics=['accuracy']
96     )
97
98     model.summary()
99
100    # Train
101    train_images, train_labels = train_data
102    val_images, val_labels = val_data
103
104    history = model.fit(
105        x=train_images,
106        y=train_labels,
107        epochs=EPOCHS,
108        batch_size=BATCH_SIZE,
109        validation_data=(val_images, val_labels),
110        callbacks=create_callbacks(model_name),
111        verbose=1
112    )
113
114    # Evaluate
115    loss, accuracy = model.evaluate(val_images, val_labels, verbose=0)
116    print(f"\n{model_name} - Final Validation Accuracy: {accuracy*100:.2f}%")
117
118    return model, history, accuracy
119
120
121 def evaluate_and_plot_metrics(model, model_name, val_images, val_labels):
122     :
123     os.makedirs(RESULTS_DIR, exist_ok=True)
124
125     # Get predictions
126     y_pred_prob = model.predict(val_images, verbose=0)
127     y_pred = np.argmax(y_pred_prob, axis=1)
128     y_true = np.argmax(val_labels, axis=1)
129
130     # Compute metrics
131     acc = accuracy_score(y_true, y_pred)
132     prec = precision_score(y_true, y_pred, average='weighted')
133
134     print(f"\n{model_name} Metrics:")
135     print(f"  Accuracy: {acc*100:.2f}%")
136     print(f"  Precision: {prec*100:.2f}%")
137
138     # Confusion matrix
139     cm = confusion_matrix(y_true, y_pred)

```

```

139
140     # Plot confusion matrix
141     plt.figure(figsize=(8, 6))
142     plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
143     plt.title(f'{model_name} Confusion Matrix')
144     plt.colorbar()
145     plt.xlabel('Predicted')
146     plt.ylabel('True')
147     plt.xticks(range(10))
148     plt.yticks(range(10))

149
150     # Add text annotations
151     for i in range(10):
152         for j in range(10):
153             plt.text(j, i, str(cm[i, j]), ha='center', va='center',
154                     color='white' if cm[i, j] > cm.max()/2 else 'black',
155                     )
156
157     plt.tight_layout()
158     plt.savefig(os.path.join(RESULTS_DIR, f'{model_name}_confusion_matrix.png'), dpi=150)
159     plt.close()
160
161     print(f" Saved: {model_name}_confusion_matrix.png")
162
163     return acc, prec
164
165 def plot_training_history(history, model_name):
166     os.makedirs(RESULTS_DIR, exist_ok=True)
167
168     # Accuracy plot
169     plt.figure(figsize=(10, 4))
170
171     plt.subplot(1, 2, 1)
172     plt.plot(history.history['accuracy'], label='Train')
173     plt.plot(history.history['val_accuracy'], label='Validation')
174     plt.title(f'{model_name} - Accuracy')
175     plt.xlabel('Epoch')
176     plt.ylabel('Accuracy')
177     plt.legend()
178     plt.grid(True)
179
180     # Loss plot
181     plt.subplot(1, 2, 2)
182     plt.plot(history.history['loss'], label='Train')
183     plt.plot(history.history['val_loss'], label='Validation')
184     plt.title(f'{model_name} - Loss')
185     plt.xlabel('Epoch')
186     plt.ylabel('Loss')
187     plt.legend()
188     plt.grid(True)
189
190     plt.tight_layout()
191     plt.savefig(os.path.join(RESULTS_DIR, f'{model_name}_training_curves.png'), dpi=150)
192     plt.close()
193
194     print(f" Saved: {model_name}_training_curves.png")
195
196
197 def get_all_models():
198     models = []
199

```

```

200 # ResNet V1
201 print("Creating ResNet V1... ")
202 try:
203     models['resnet_v1'] = get_resnetv1(
204         num_classes=NUM_CLASSES,
205         input_shape=DATA_SHAPE,
206         depth=8,
207         dropout=0.2
208     )
209 except Exception as e:
210     print(f"Error creating ResNet: {e}")
211
212 # SqueezeNet V1.1
213 print("Creating SqueezeNet V1.1... ")
214 try:
215     models['squeezenet_v11'] = get_squeezenetv11(
216         num_classes=NUM_CLASSES,
217         input_shape=DATA_SHAPE,
218         dropout=0.5
219     )
220 except Exception as e:
221     print(f"Error creating SqueezeNet: {e}")
222
223 # EfficientNet V2-B0
224 print("Creating EfficientNet V2-B0... ")
225 try:
226     models['efficientnet_v2_b0'] = get_efficientnetv2(
227         input_shape=DATA_SHAPE,
228         model_type='B0',
229         num_classes=NUM_CLASSES,
230         dropout=0.2,
231         pretrained_weights=None # Train from scratch
232     )
233 except Exception as e:
234     print(f"Error creating EfficientNet: {e}")
235
236 # MobileNetV1
237 print("Creating MobileNetV1... ")
238 try:
239     models['mobilenet_v1'] = get_mobilenetv1(
240         input_shape=DATA_SHAPE,
241         alpha=0.25,
242         num_classes=NUM_CLASSES,
243         dropout=0.2,
244         pretrained_weights=None
245     )
246 except Exception as e:
247     print(f"Error creating MobileNetV1: {e}")
248
249 # MobileNetV2
250 print("Creating MobileNetV2... ")
251 try:
252     models['mobilenet_v2'] = get_mobilenetv2(
253         input_shape=DATA_SHAPE,
254         alpha=0.35,
255         num_classes=NUM_CLASSES,
256         dropout=0.2,
257         pretrained_weights=None
258     )
259 except Exception as e:
260     print(f"Error creating MobileNetV2: {e}")
261
262 # FDMobileNet
263 print("Creating FDMobileNet... ")

```

```

264     try:
265         models['fdmobilenet'] = get_fdmobilenet(
266             input_shape=DATA_SHAPE,
267             num_classes=NUM_CLASSES,
268             alpha=0.25,
269             dropout=0.2
270         )
271     except Exception as e:
272         print(f"Error creating FDMobileNet: {e}")
273
274 # ST EfficientNet LC V1
275 print("Creating ST EfficientNet LC V1...")
276 try:
277     models['st_efficientnet_lc_v1'] = get_st_efficientnet_lc_v1(
278         input_shape=DATA_SHAPE,
279         num_classes=NUM_CLASSES,
280         dropout=0.2
281     )
282 except Exception as e:
283     print(f"Error creating ST EfficientNet LC: {e}")
284
285 # ST FDMobileNet V1
286 print("Creating ST FDMobileNet V1...")
287 try:
288     models['st_fdmobilenet_v1'] = get_st_fdmobilenet_v1(
289         input_shape=DATA_SHAPE,
290         num_classes=NUM_CLASSES,
291         dropout=0.2
292     )
293 except Exception as e:
294     print(f"Error creating ST FDMobileNet: {e}")
295
296 return models
297
298
299 def main():
300     os.makedirs(MODELS_DIR, exist_ok=True)
301     train_data, val_data = load_mnist_data()
302
303     # Create models
304     print("\n" + "="*60)
305     print("Creating CNN Models")
306     print("=".*60)
307     models = get_all_models()
308
309     # Train each model
310     results = {}
311     val_images, val_labels = val_data
312
313     for model_name, model in models.items():
314         try:
315             trained_model, history, accuracy = train_model(model,
316                 model_name, train_data, val_data)
317
318             # Save training curves
319             plot_training_history(history, model_name)
320
321             # Compute metrics and save confusion matrix
322             acc, prec = evaluate_and_plot_metrics(trained_model,
323                 model_name, val_images, val_labels)
324             results[model_name] = {'accuracy': acc, 'precision': prec}
325
326             # Clear session to free memory
327             keras.backend.clear_session()

```

```

326
327     except Exception as e:
328         print(f"Error training {model_name}: {e}")
329         results[model_name] = {'accuracy': 0.0, 'precision': 0.0}
330
331 # Print final results
332 print("\n" + "*60)
333 print("TRAINING RESULTS SUMMARY")
334 print("*60)
335 print(f"{'Model':<25} {'Accuracy':>10} {'Precision':>10}")
336 print("-*47)
337 for model_name, metrics in sorted(results.items(), key=lambda x: x
338 [1]['accuracy'], reverse=True):
339     print(f"{model_name:<25} {metrics['accuracy']*100:>9.2f}% {metrics['precision']*100:>9.2f}%")
340
341 print(f"\nModels saved to: {MODELS_DIR}")
342 return results
343
344 if __name__ == "__main__":
345     main()

```

Quantization and Conversion: Initially, we attempted to convert the models directly using y5 to tflite workflow. However, we found that without quantization, the model sizes were too large to fit into the limited Flash memory of the STM32F446RE. Therefore, we implemented a quantization script that converts the Keras models (.h5) to TensorFlow Lite (.tflite) format while applying full integer quantization ().

The quantization process uses a representative dataset drawn from the MNIST training data to estimate the dynamic range of activations. This ensures that the conversion from floating-point to integer does not significantly degrade accuracy.

Listing 2.2: Quantization script (quantize_models.py)

```

1 import os
2 import sys
3 import numpy as np
4 import tensorflow as tf
5 from tensorflow import keras
6
7 # Configuration
8 MODELS_DIR = os.path.join(os.path.dirname(__file__), 'models')
9 DATA_SHAPE = (32, 32, 3)
10
11 def prepare_tensor(images, out_shape):
12     images = tf.expand_dims(images, axis=-1)
13     images = tf.repeat(images, 3, axis=-1)
14     images = tf.image.resize(images, out_shape[:2])
15     images = images / 255.0
16     return images
17
18 def get_representative_dataset():
19     print("Loading MNIST for representative dataset...")
20     (train_images, _), _ = tf.keras.datasets.mnist.load_data()
21
22     # Preprocess a subset of images

```

```

23 # We use 100 samples as recommended for representative datasets
24 num_calibration_steps = 100
25
26 # Shuffle and pick a subset
27 indices = np.random.choice(len(train_images), num_calibration_steps,
28     replace=False)
28 calibration_images = train_images[indices]
29
30 # Preprocess
31 print("Preprocessing representative data...")
32 calibration_images = prepare_tensor(calibration_images, DATA_SHAPE)
33
34 def representative_data_gen():
35     for input_value in calibration_images:
36         # Model expects [1, 32, 32, 3]
37         input_value = tf.expand_dims(input_value, axis=0)
38         yield [input_value]
39
40 return representative_data_gen
41
42 def quantize_model(h5_path):
43     base_name = os.path.splitext(os.path.basename(h5_path))[0]
44     output_path = os.path.join(os.path.dirname(h5_path), f'{base_name}_quant.tflite')
45
46     try:
47         print(f"\nProcessing: {base_name}")
48
49         # Load Keras model
50         model = keras.models.load_model(h5_path)
51
52         # Create converter
53         converter = tf.lite.TFLiteConverter.from_keras_model(model)
54
55         # Set optimization flags
56         converter.optimizations = [tf.lite.Optimize.DEFAULT]
57
58         # Set representative dataset
59         converter.representative_dataset = get_representative_dataset()
60
61         # Ensure full integer quantization
62         converter.target_spec.supported_ops = [tf.lite.OpsSet.TFLITE_BUILTINS_INT8]
63
64         # Convert
65         print("Converting with quantization...")
66         tflite_model = converter.convert()
67
68         # Save
69         with open(output_path, 'wb') as f:
70             f.write(tflite_model)
71
72         original_size = os.path.getsize(h5_path) / 1024
73         quant_size = len(tflite_model) / 1024
74
75         print(f"Saved: {output_path}")
76         print(f"Original H5 size: {original_size:.1f} KB")
77         print(f"Quantized size: {quant_size:.1f} KB")
78         print(f"Reduction ratio: {original_size/quant_size:.2f}x")
79
80         return {
81             'name': base_name,
82             'original_size': original_size,
83             'quant_size': quant_size

```

```

84     }
85
86     except Exception as e:
87         print(f"Error converting {base_name}: {e}")
88         return None
89
90 def main():
91     if not os.path.exists(MODELS_DIR):
92         print(f"Models directory not found: {MODELS_DIR}")
93         return
94
95     h5_files = [f for f in os.listdir(MODELS_DIR) if f.endswith('.h5')]
96
97     if not h5_files:
98         print("No .h5 files found to quantize.")
99         return
100
101    results = []
102    for f in h5_files:
103        h5_path = os.path.join(MODELS_DIR, f)
104        res = quantize_model(h5_path)
105        if res:
106            results.append(res)
107
108 # Summary
109 print("\n" + "="*60)
110 print("QUANTIZATION SUMMARY")
111 print("="*60)
112 print(f"{'Model':<25} {'Orig (KB)':>10} {'Quant (KB)':>12} {'Ratio':>8}")
113 print("-"*57)
114
115 for r in results:
116     ratio = r['original_size'] / r['quant_size']
117     print(f"{r['name']:<25} {r['original_size']:>10.1f} {r['quant_size']:>12.1f} {ratio:>8.2f}x")
118
119 if __name__ == "__main__":
120     main()

```

Step 2: ST Edge AI Developer Cloud Optimization

After generating the quantized .tflite models, we utilized the **ST Edge AI Developer Cloud** [2] to validate their deployability on the STM32F446RE. This platform automates the complex task of mapping neural network operators to the specific hardware accelerator and memory architecture of the target MCU.

Model Analysis and Validation: We uploaded the quantized models to the cloud platform and selected the **NUCLEO-F446RE** as the target board. The service analyzes the model's structure to determine the required Flash memory (for weights) and RAM (for activation buffers).




Figure 2.1: ST Edge AI Developer Cloud analysis dashboard showing SqueezeNetV11 memory usage exceeds flash size.

The analysis revealed a critical hardware limitation. Despite having sufficient RAM for the activation buffers, the binary size for both **EfficientNet** and **SqueezeNet** exceeded the 512kB Flash memory limit of the STM32F446RE. As a result, these models could not be deployed.

Project Generation: The remaining models **ResNetV1**, **MobileNetV1**, **FD-MobileNetV1**, and **ST-FD-MobileNetV1** successfully passed the memory validation checks. For these models, we ran the cloud benchmark to estimate inference latency and then generated the optimized C-code. The platform provided a downloadable STM32CubeIDE project ('.zip') containing the X-CUBE-AI library pre-configured with the network weights and runtime environment.




Figure 2.2: Successful generation of the STM32CubeIDE project for ResNetV1.

Step 3: STM32 Firmware Implementation

The generated STM32CubeIDE projects contain the X-CUBE-AI library, which provides the runtime environment for the neural networks. However, the default `main.c` file requires modification to handle real-time data ingestion and specific input preprocessing requirements.

We implemented a unified firmware structure that works across all deployed models (ResNetV1, MobileNetV1, etc.). The core logic involves three key components: Preprocessing, Initialization, and the Inference Loop.

Input Preprocessing: The models were trained on 32x32 RGB images, but the incoming data from the UART interface is raw 28x28 grayscale (MNIST format). We implemented a `PrepareInput` function to bridge this gap. This function performs:

1. **Resizing:** Bilinear interpolation from 28x28 to 32x32 pixels.

2. **Normalization:** Scaling pixel values from [0, 255] to [0.0, 1.0].

3. **Channel Expansion:** Replicating the single grayscale channel into 3 RGB channels to match the network's input tensor shape (32, 32, 3).

Listing 2.3: Preprocessing function handling resize and normalization

```
1  /* Bilinear interpolation resize 28x28 -> 32x32, grayscale -> RGB,
2   * normalize */
3
4  static void PrepareInput(uint8_t *img28, float *in32) {
5      int x32, y32, c;
6      float scale = 28.0f / 32.0f;
7
8      for (y32 = 0; y32 < IMG_SIZE_32; y32++) {
9          for (x32 = 0; x32 < IMG_SIZE_32; x32++) {
10             /* Map 32x32 coord to 28x28 */
11             float src_x = x32 * scale;
12             float src_y = y32 * scale;
13
14             int x0 = (int)src_x;
15             int y0 = (int)src_y;
16             int x1 = (x0 < IMG_SIZE_28 - 1) ? x0 + 1 : x0;
17             int y1 = (y0 < IMG_SIZE_28 - 1) ? y0 + 1 : y0;
18
19             float dx = src_x - x0;
20             float dy = src_y - y0;
21
22             /* Bilinear interpolation */
23             float p00 = img28[y0 * IMG_SIZE_28 + x0];
24             float p10 = img28[y0 * IMG_SIZE_28 + x1];
25             float p01 = img28[y1 * IMG_SIZE_28 + x0];
26             float p11 = img28[y1 * IMG_SIZE_28 + x1];
27
28             float val = p00 * (1 - dx) * (1 - dy) +
29                         p10 * dx * (1 - dy) +
30                         p01 * (1 - dx) * dy +
31                         p11 * dx * dy;
32
33             /* Normalize to [0,1] */
34             float norm_val = val / 255.0f;
35
36             /* Replicate to RGB channels (HWC format: 32x32x3) */
37             int idx = (y32 * IMG_SIZE_32 + x32) * 3;
38             for (c = 0; c < 3; c++) {
39                 in32[idx + c] = norm_val;
40             }
41         }
42     }
43 }
```

Network Initialization: The AI_Init function configures the neural network context. It links the model weights (stored in Flash) and the activation buffers (allocated in RAM) to the inference engine.

Listing 2.4: Network initialization using STAI API

```
1  static int AI_Init(void) {
```

```

2   if (stai_network_init((stai_network *)net_ctx) != STAI_SUCCESS)
3     return -1;
4
5   /* Set memory for activations */
6   if (stai_network_set_activations((stai_network *)net_ctx, act_ptrs, 1)
7     != STAI_SUCCESS)
8     return -1;
9
10  /* Link input and output buffers */
11  if (stai_network_set_inputs((stai_network *)net_ctx, in_ptrs, 1) !=
12    STAI_SUCCESS)
13    return -1;
14
15  if (stai_network_set_outputs((stai_network *)net_ctx, out_ptrs, 1) !=
16    STAI_SUCCESS)
17    return -1;
18
19  return 0;
20 }
```

Main Inference Loop: The main loop utilizes a blocking UART receive call to wait for a synchronization byte (0xBB) and the image payload (784 bytes). Upon reception, the image is processed, inference is executed synchronously, and the predicted class with the highest probability (ArgMax) is transmitted back to the PC.

Listing 2.5: Main execution loop

```

1  /* ... Initialization code ... */
2
3  while (1) {
4    /* Wait for Sync Byte */
5    if (HAL_UART_Receive(&huart2, &sync, 1, HAL_MAX_DELAY) == HAL_OK &&
6      sync == SYNC_BYTEx {
7      HAL_GPIO_WritePin(LD2_GPIO_Port, LD2_Pin, GPIO_PIN_RESET); // LED
8      OFF = Busy
9
10   /* Receive Raw Image */
11   if (HAL_UART_Receive(&huart2, img_buf, IMG_BYTES, 5000) == HAL_OK)
12     {
13
14     /* Preprocess: 28x28 -> 32x32x3 float */
15     PrepareInput(img_buf, in_buf);
16
17     /* Run Inference */
18     if (stai_network_run((stai_network *)net_ctx, STAI_MODE_SYNC) ==
19       STAI_SUCCESS) {
20       res = ArgMax(out_buf, 10); /* Get Class Index */
21     } else {
22       res = 0xFF; /* Error */
23     }
24
25     /* Transmit Result */
26     HAL_UART_Transmit(&huart2, &res, 1, 100);
27   }
28   HAL_GPIO_WritePin(LD2_GPIO_Port, LD2_Pin, GPIO_PIN_SET); // LED ON
29   = Ready
30 }
```

Step 4: Python UART Client

To validate the deployed model, we developed a Python script ('mnist_uart.py') that acts as the interface between the PC and the STM32F446RE. This script handles image loading, preprocessing (resizing to 28x28, inverting colors if necessary for white-on-black digits), and UART communication.

The client sends a synchronization byte ('0xBB') followed by the raw pixel data (784 bytes) to the board. It then waits for the predicted digit index to be returned.

Listing 2.6: Python UART client for MNIST inference (mnist_uart.py)

```
1 import argparse
2 import numpy as np
3 import cv2
4 import serial
5
6 SYNC_BYTE = 0xBB
7
8 def load_image(path):
9     img = cv2.imread(path, cv2.IMREAD_GRAYSCALE)
10    if img is None:
11        raise ValueError(f"Cannot load: {path}")
12    img = cv2.resize(img, (28, 28))
13    # Invert colors if the image is black-on-white (like standard
14    # writing)
15    # MNIST requires white-on-black
16    if np.mean(img) > 127:
17        img = 255 - img
18    return img
19
20 def send_and_receive(img, port='COM9', baudrate=115200):
21     try:
22         with serial.Serial(port, baudrate, timeout=5) as ser:
23             ser.write(bytes([SYNC_BYTE]))
24             ser.write(img.flatten().tobytes())
25             print(f"Sent 784 bytes")
26             result = ser.read(1)
27             return result[0] if len(result) == 1 else -1
28     except Exception as e:
29         print(f"Error: {e}")
30         return -1
31
32 def display(digit, img=None):
33     canvas = np.zeros((600, 800, 3), dtype=np.uint8)
34     cv2.putText(canvas, str(digit), (280, 400), cv2.FONT_HERSHEY_SIMPLEX,
35                 15, (0, 255, 0), 30)
36     cv2.putText(canvas, "Recognized Digit", (150, 60), cv2.
37                 FONT_HERSHEY_SIMPLEX, 1.5, (255, 255, 255), 3)
38     if img is not None:
39         # Show the input image in the corner
40         canvas[80:220, 20:160] = cv2.cvtColor(
41             cv2.resize(img, (140, 140), interpolation=cv2.INTER_NEAREST),
42             cv2.COLOR_GRAY2BGR
43         )
44     cv2.imshow('Result', canvas)
45     cv2.waitKey(0)
46     cv2.destroyAllWindows()
```

```

44
45 def main():
46     parser = argparse.ArgumentParser()
47     parser.add_argument('image', help='Path to the image file')
48     parser.add_argument('--port', default='COM9', help='Serial port (
        default: COM9)')
49     args = parser.parse_args()
50
51     img = load_image(args.image)
52     print(f"Sending 28x28 image to MCU...")
53     digit = send_and_receive(img, args.port)
54
55     if 0 <= digit <= 9:
56         print(f"*** Recognized: {digit} ***")
57         display(digit, img)
58     else:
59         print("Failed or Timed Out")
60
61 if __name__ == '__main__':
62     main()

```

Usage: The script is executed from the command line. It requires the path to the image file and optionally accepts a specific COM port. The recognized digit is displayed in an OpenCV window along with a visual confirmation.

Example command:

```
python mnist_uart.py two.png --port COM9
```

Results

Model Performance Analysis

All six trained models achieved high convergence on the MNIST dataset, demonstrating validation accuracies exceeding 95%. Figure 2.3 illustrates the training progression over 20 epochs. ResNetV1 and MobileNetV1 showed the most stable convergence, reaching optimal accuracy within the first 5 epochs.




Figure 2.3: Training accuracy and loss curves for all evaluated models.

The confusion matrices in Figure 2.4 provide a deeper insight into classification performance. The diagonal dominance across all models confirms robust feature extraction. **ResNetV1** and **EfficientNetV2** exhibited the fewest off-diagonal errors, indicating superior precision. But the thing is SqueezeNet has some off labels as can be seen in the figure. That may be why we couldn't be able to fit into MCU because of some computational error that happened during the training.




Figure 2.4: Confusion matrices comparing the classification performance of the six architectures.

Hardware Resource Utilization

The critical constraint for this project was the 512kB Flash memory of the STM32F446RE. Table 2.1 summarizes the memory footprint of each model after quantization and optimization by the ST Edge AI Cloud.

Table 2.1: Resource Utilization and Deployment Status on STM32F446RE


Model Architecture	Flash Usage (kB)	RAM Usage (kB)	Deployment Status
ResNetV1	112	44	Success
FDMobileNetV1	160	13	Success
ST-FDMobileNetV1	178	13	Success
MobileNetV1	252	13	Success
EfficientNetV2-B0	6,000	84	Failed (Flash Overflow)
SqueezeNetV1.1	767	17	Failed (Flash Overflow)

Despite their high accuracy, both EfficientNetV2 (6MB) and SqueezeNetV1.1 (767kB) exceeded the available Flash memory. While SqueezeNet is often cited as a "lightweight"


architecture, its parameter density proved too high for this specific microcontroller class without external memory. Conversely, ResNetV1 was the most balanced candidate, consuming only 112kB of Flash while maintaining high accuracy.

Real-Time Inference Verification

The four compatible models were successfully deployed to the MCU. We verified real-time performance using the Python UART client. Figure 2.5 demonstrates the successful classification of handwritten digits sent from the PC to the STM32F446RE.



(a) ResNetV1 correctly classifying digit '2'.



(b) ST-FDMobileNetV1 correctly classifying digit '5'.

Figure 2.5: Real-time inference results from the STM32F446RE visualized via the Python client and showing computed probabilities of the label shown in STM32CubeIDE

3. CONCLUSION

In this homework, we successfully deployed deep CNN architectures onto the STM32F446RE MCU for handwritten digit recognition. The main goal was to evaluate different neural network topologies and implement a complete edge AI pipeline, navigating the strict memory constraints of a mid-range microcontroller.

We trained multiple architectures, including ResNetV1, MobileNetV1, SqueezeNet, and EfficientNet, using the MNIST dataset. Unlike previous assignments that relied on manual feature extraction like Hu Invariant Moments (Hu Moments), this work utilized deep learning models that process raw image data directly. We developed a Python client to transmit images via Universal Asynchronous Receiver Transmitter (UART), and implemented firmware on the MCU to perform real-time preprocessing (resizing and normalization) and inference.

A critical insight from this work was the impact of model complexity on hardware feasibility. While models like EfficientNetV2 and SqueezeNet offer high theoretical performance, we observed that they exceeded the 512kB Flash memory limit of the STM32F446RE, rendering them undeployable without external memory. Conversely, **ResNetV1** and **MobileNetV1**, combined with INT8 quantization, provided an optimal balance, fitting comfortably within the hardware resources while maintaining high classification accuracy.

We used the **ST Edge AI Developer Cloud** to bridge the gap between TensorFlow and the embedded C environment. This tool, along with our custom quantization scripts, allowed us to optimize model footprints significantly. This assignment highlighted the importance of hardware aware design, shows that successful edge AI deployment requires not just training accurate models but strictly adjusting to the physical constraints of the target device.

BIBLIOGRAPHY

- [1] C. Ünsalan, B. Höke, E. Atmaca, *Embedded Machine Learning with Microcontrollers: Applications on STM32 Boards*, Springer Nature, 2025.
- [2] STMicroelectronics, St edge ai developer cloud, <https://stedgeai-dc.st.com> (2024).
- [3] M. Abadi, et al., Tensorflow: Large-scale machine learning on heterogeneous systems, <https://www.tensorflow.org> (2015).